



Programming an Autonomous Driving Car with J.K. Rollin

Jitendra Awasthi & Kam Mirhosseini

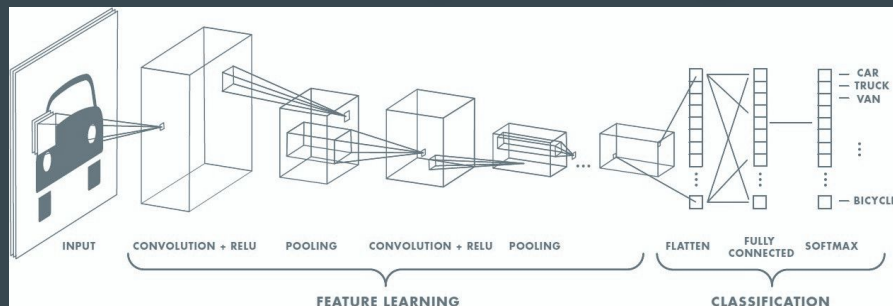
Introduction: Autonomous cars

What is an autonomous car?

An autonomous car is a vehicle capable of sensing the environment around it and operating without human involvement.

- Reduce traffic congestion
- Cut transportation costs (in terms of vehicles, fuel, and infrastructure)
- Reduce traffic accidents
- Free up parking lots for other uses
- Reduce urban CO2 emissions worldwide

Introduction: Deep Learning & CNNs



Before Deep Learning Era

- HAAR (Viola & Jones, 2001)
- Local Binary Patterns (LBP)(Ojala et al., 1996).

Rise of Deep Learning and CNN

- Recent development in hardware and affordability
- Increase of computational power
- Availability of big data

CNNs were some of the first deep models to solve commercial applications.

One of the first examples of these networks is the AT&T in 1990.

The ImageNet challenge led to some amazing computer vision networks such as, AlexNet and ResNet.

Introduction: The task and approaches

The task: Predicting angle and speed based off of images

Main approaches:

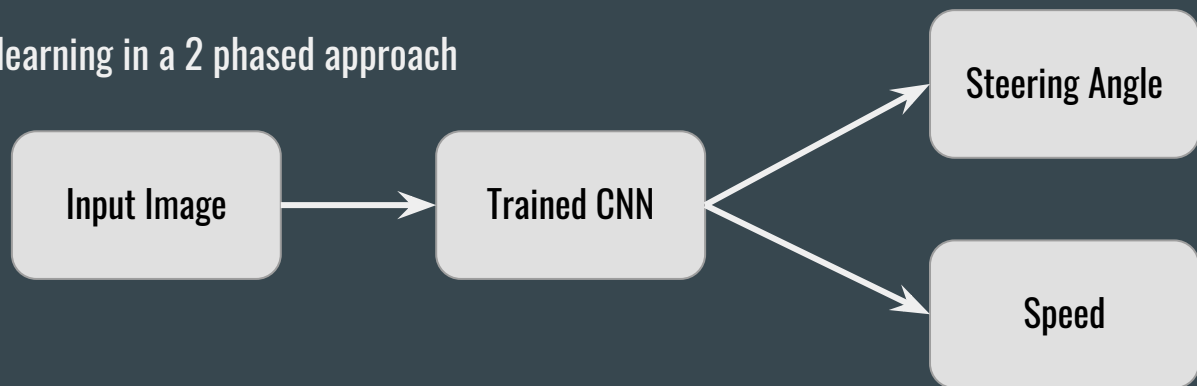
- Rule-based learning
- End-to-end learning

Choice of algorithms:

- Supervised learning
- Unsupervised learning
- Reinforcement Learning



We used end-to-end supervised learning in a 2 phased approach

- Object classification
- Regression



Data

1. Total images: 13793
2. We observed that left turn images are under-represented.
3. Left arrows were less than right arrows
4. Tricky images e.g. tree on the side and special scenarios
5. Gender bias towards female figures
6. Inappropriate values in training_norm.csv:
 - Filename 4725 speed should be 0
 - Filename 1447 angle should be close 0.5



Data Processing: Augmentation

- Noise
- Brightness
- Exposure
- Blur
- Zoom
- Pan
- Flip
- Rotation

We had to flip the angle accordingly.

As there was more right turns on the training data, the threshold for applying random flip augmentation was increased to 0.7 in order to have even number of left/right turns.





Models

Tools used in our models

APIs:

- Tensorflow Keras functional API
- For the Nvidia model, sequential API

Layers:

- CNNs
- Dense
- Dropout (Srivastava et al., 2014)
- Batch Normalisation

Activation function: ELU and Sigmoid

Optimizer:

- The Adam optimizer (Kingma & Ba, 2014)

Regularizer:

- Tried regularisation but we did not get any improvement

Early stopping in label detection with a set patience

Checkpoint callbacks to be able to go back to a previous model if overfitting occurs.

Model 1: Nvidia Model (Bojarski et al., 2016)

Input and preprocessing: Full image resized to half

Before trying to create our own model we tried out the Nvidia Model.

We changed the output dense layer to give 2 predictions instead of 1, which gave us our first kaggle submission.


Final Loss: 0.07582

Score: 0.05893

Score: 0.04830

| Layer (type) | Output Shape | Param # |
|---------------------------|--------------------|---------|
| conv2d_5 (Conv2D) | (None, 31, 98, 24) | 1824 |
| conv2d_6 (Conv2D) | (None, 14, 47, 36) | 21636 |
| conv2d_7 (Conv2D) | (None, 5, 22, 48) | 43248 |
| conv2d_8 (Conv2D) | (None, 3, 20, 64) | 27712 |
| dropout_2 (Dropout) | (None, 3, 20, 64) | 0 |
| conv2d_9 (Conv2D) | (None, 1, 18, 64) | 36928 |
| flatten_1 (Flatten) | (None, 1152) | 0 |
| dropout_3 (Dropout) | (None, 1152) | 0 |
| dense_4 (Dense) | (None, 100) | 115300 |
| dense_5 (Dense) | (None, 50) | 5050 |
| dense_6 (Dense) | (None, 10) | 510 |
| dense_7 (Dense) | (None, 2) | 22 |
| Total params: 252,230 | | |
| Trainable params: 252,230 | | |
| Non-trainable params: 0 | | |
| None | | |

Transfer Learning



We removed the last layers of each model and concatenate the last layers and added extra dense layers as needed.

We froze the entire model except the new added layers and we trained the new layers for a few epochs to warm up the new layers.

We then unfroze the entire model and trained with a low learning rate.


We tried a combination of training approaches.

Label Detector Model Based on Nvidia

We used LabelIMG and Roboflow to label the objects and make augmentations.

Our labels consisted of:

1. Person
2. Red Light
3. Green Light
4. Right Sign
5. Left Sign
6. Tree
7. Box
8. Car




| Layer (type) | Output Shape | Param # |
|---------------------------|--------------------|---------|
| conv2d (Conv2D) | (None, 58, 78, 24) | 1824 |
| conv2d_1 (Conv2D) | (None, 27, 37, 36) | 21636 |
| conv2d_2 (Conv2D) | (None, 12, 17, 48) | 43248 |
| conv2d_3 (Conv2D) | (None, 10, 15, 64) | 27712 |
| dropout (Dropout) | (None, 10, 15, 64) | 0 |
| conv2d_4 (Conv2D) | (None, 8, 13, 64) | 36928 |
| flatten (Flatten) | (None, 6656) | 0 |
| dropout_1 (Dropout) | (None, 6656) | 0 |
| dense (Dense) | (None, 100) | 665700 |
| dense_1 (Dense) | (None, 50) | 5050 |
| dropout_2 (Dropout) | (None, 50) | 0 |
| dense_2 (Dense) | (None, 10) | 510 |
| dense_3 (Dense) | (None, 8) | 88 |
| Total params: 802,696 | | |
| Trainable params: 802,696 | | |
| Non-trainable params: 0 | | |
| None | | |

Model 2: The Twin Engine Model

Transfer Learning using Label Detector for speed

- Full image
- Resized



Label Detector Model Based on Nvidia

Conv
24

Conv
36

Conv
48

Conv
64

Conv
64

Flatten

Dropout

Concatenate

Dense
100

Dense
50

Dense
10

Dense 2

Model for angle detection

- Half Image
- Resized
- Grayscale

Kaggle Score: 0.03358

Model 2: Experiment

Label Detector Model Based on Nvidia

Dense 50

Dense 10

Dense 1
Speed

Conv 24

Conv 36

Conv 48

Conv 64

Conv 64

Flatten

Dropout

Concatenate

Dense
100

Dense 50

Dense 10

Dense 1
Angle

Single Engine Model

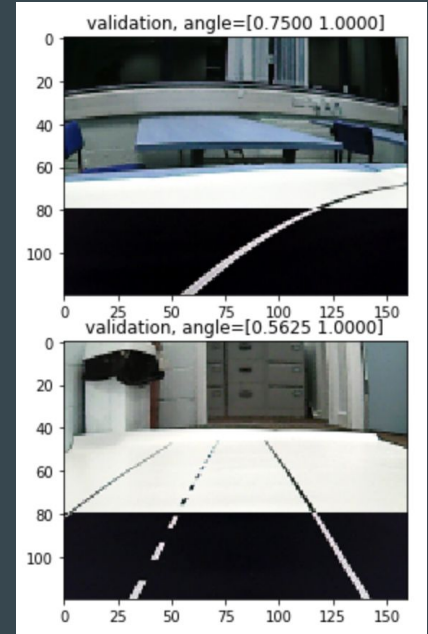
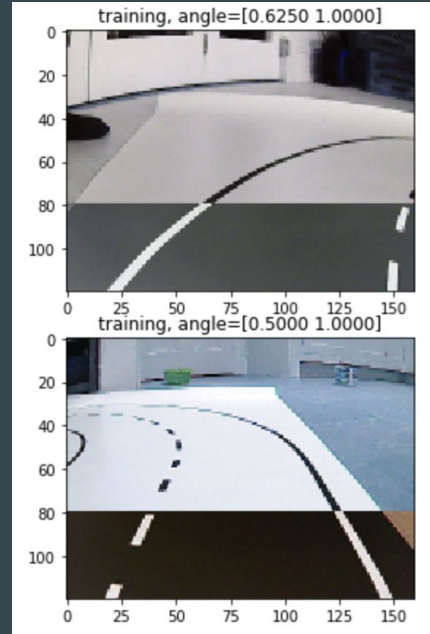
We were giving two sets of images to the twin engine model so we decided to get the same effect from a single image.

At this stage, we decided to try segmentation.

We inverted the lower segment of the images hoping it would learn both angle and objects in a single model.

Once again, the angle was not predicted based on the arrows.

However, we came back to this model later



We moved on from the Nvidia model to our own model

Divide and Rule (PREDICT)

Angle:

Full Image

Preprocessing:

- Grayscale
- Inverted lower $\frac{1}{3}$ of the image
- Resize & Normalise

Label Detection

Transfer Learning

Speed:

Full Image

Preprocessing:

- Resize & Normalise

Label Detection

Transfer Learning

Model 3: Twin Engine v.2.0

Phase 1: Label Detection for Speed related objects

For Speed we used labels:

1. Person
2. Red Light
3. Green Light
4. Tree
5. Box
6. Car

| Layer (type) | Output Shape | Param # |
|---|-----------------------|---------|
| input_1 (InputLayer) | [(None, 120, 160, 3)] | 0 |
| conv2d (Conv2D) | (None, 118, 158, 32) | 896 |
| conv2d_1 (Conv2D) | (None, 116, 156, 32) | 9248 |
| max_pooling2d (MaxPooling2D) | (None, 58, 78, 32) | 0 |
| conv2d_2 (Conv2D) | (None, 56, 76, 48) | 13872 |
| conv2d_3 (Conv2D) | (None, 54, 74, 48) | 20784 |
| max_pooling2d_1 (MaxPooling2D) | (None, 27, 37, 48) | 0 |
| conv2d_4 (Conv2D) | (None, 25, 35, 64) | 27712 |
| conv2d_5 (Conv2D) | (None, 23, 33, 64) | 36928 |
| max_pooling2d_2 (MaxPooling2D) | (None, 11, 16, 64) | 0 |
| flatten (Flatten) | (None, 11264) | 0 |
| dropout (Dropout) | (None, 11264) | 0 |
| dense (Dense) | (None, 256) | 2883840 |
| dense_1 (Dense) | (None, 100) | 25700 |
| batch_normalization (Batch Normalization) | (None, 100) | 400 |
| dense_2 (Dense) | (None, 50) | 5050 |
| dense_3 (Dense) | (None, 10) | 510 |
| angle (Dense) | (None, 1) | 11 |
| Total params: 3,024,951 | | |
| Trainable params: 3,024,751 | | |
| Non-trainable params: 200 | | |

Model 3: Twin Engine v.2.0

Phase 1: Label Detection for Angle related objects


For Angle we used:

1. Right Sign
2. Left sign

We also experimented labelling the road to detect angles. Starting with 0, 0.5 and 1



0



0.5



1

| Layer (type) | Output Shape | Param # |
|---|-----------------------|---------|
| input_1 (InputLayer) | [(None, 120, 160, 1)] | 0 |
| conv2d (Conv2D) | (None, 118, 158, 32) | 320 |
| conv2d_1 (Conv2D) | (None, 116, 156, 32) | 9248 |
| max_pooling2d (MaxPooling2D) | (None, 58, 78, 32) | 0 |
| conv2d_2 (Conv2D) | (None, 56, 76, 48) | 13872 |
| conv2d_3 (Conv2D) | (None, 54, 74, 48) | 20784 |
| max_pooling2d_1 (MaxPooling2D) | (None, 27, 37, 48) | 0 |
| conv2d_4 (Conv2D) | (None, 25, 35, 64) | 27712 |
| conv2d_5 (Conv2D) | (None, 23, 33, 64) | 36928 |
| max_pooling2d_2 (MaxPooling2D) | (None, 11, 16, 64) | 0 |
| Flatten (Flatten) | (None, 11264) | 0 |
| dropout (Dropout) | (None, 11264) | 0 |
| dense (Dense) | (None, 256) | 2883840 |
| dense_1 (Dense) | (None, 100) | 25700 |
| batch_normalization (Batch Normalization) | (None, 100) | 400 |
| dense_2 (Dense) | (None, 50) | 5050 |
| dense_3 (Dense) | (None, 10) | 510 |
| angle (Dense) | (None, 1) | 11 |
| Total params: 3,024,375 | | |
| Trainable params: 3,024,175 | | |
| Non-trainable params: 200 | | |

Model 3: Twin Engine v.2.0

Phase 2: Transfer Learning

In phase 2 we trained to predict for angle and speed and combined them using all the steps for Transfer learning.

This model was more successful and resulted in our best kaggle score of 0.03197

Why we did not use this was because the model was big in terms of inference time.

We tried many approaches to cut down the model.

| Layer (type) | Output Shape | Param # |
|---|-----------------------|---------|
| input_1 (InputLayer) | [(None, 120, 160, 1)] | 0 |
| conv2d (Conv2D) | (None, 118, 158, 32) | 320 |
| conv2d_1 (Conv2D) | (None, 116, 156, 32) | 9248 |
| max_pooling2d (MaxPooling2D) | (None, 58, 78, 32) | 0 |
| conv2d_2 (Conv2D) | (None, 56, 76, 48) | 13872 |
| conv2d_3 (Conv2D) | (None, 54, 74, 48) | 20784 |
| max_pooling2d_1 (MaxPooling2D) | (None, 27, 37, 48) | 0 |
| conv2d_4 (Conv2D) | (None, 25, 35, 64) | 27712 |
| conv2d_5 (Conv2D) | (None, 23, 33, 64) | 36928 |
| max_pooling2d_2 (MaxPooling2D) | (None, 11, 16, 64) | 0 |
| flatten (Flatten) | (None, 11264) | 0 |
| dropout (Dropout) | (None, 11264) | 0 |
| dense (Dense) | (None, 256) | 2883840 |
| dense_1 (Dense) | (None, 100) | 25700 |
| batch_normalization (Batch Normalization) | (None, 100) | 400 |
| dense_2 (Dense) | (None, 50) | 5050 |
| dense_3 (Dense) | (None, 10) | 510 |
| angle (Dense) | (None, 1) | 11 |
| Total params: 3,024,375 | | |
| Trainable params: 3,024,175 | | |
| Non-trainable params: 200 | | |

Twin Engine v2.1

The layer we removed was *dense(256)* layer with 2883840 number of parameters.

This model started to perform faster, however, integrating this model with the code we realised we still need to further reduce the inference time and the kaggle score went down.

Kaggle Score: 0.04030

Twin Engine v2.2

We modified the first two convolutional layers by changing the *filter size (5, 5)* and added stride (2, 2). We added the *dense(256)* layer back.




The Final Day

The kaggle submission for the Twin Engine model that used labels, was 0.06192

We realised the new images added recently were distorted by over-augmentation to a point where different arrows and traffic lights could not be distinguished.

We removed the overly augmented images and trained overnight on the remaining data without training for object detection.

In the morning we made a better representation of all special scenarios into one folder. We re-trained the model for speed and angle separately and combined them. We saw improvement on some special scenarios but we did not realise this training resulted in some association between straight road and (speed=0).



How we would do the project if we had to do it again

Begin the project by focusing on the dataset

Make sure every scenario is properly represented

To have better control on training we would use a custom data generator.

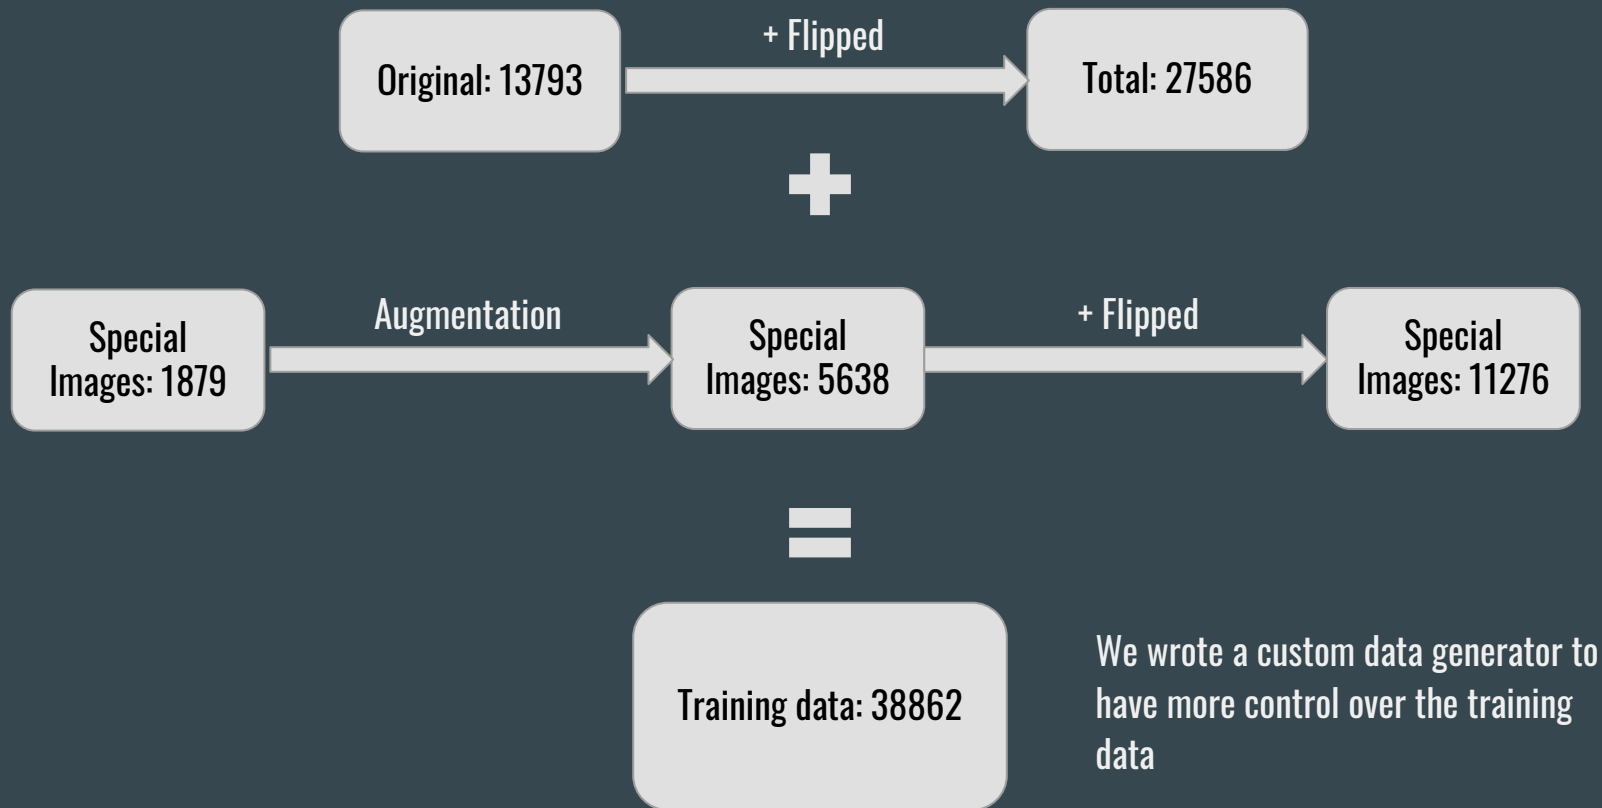
Retraining the model on special subset is tricky and needs more attention.

Labelling the obstacles as 1 category (may or may not help). Try anyways!

We learnt a lot by trying various approaches, however, if we focused on one model and identified the problems we would have had a better performance.

Too much excitement over kaggle :)



Redesigning the training data



Revisiting The Sweet & Simple Model

Input and Preprocessing:


- Full Image
- Inverted the lower (Approx.) $\frac{1}{3}$ of the image
- Resized to half
- Normalised



| Layer (type) | Output Shape | Param # |
|---|-----------------------|---------|
| input_1 (InputLayer) | [(None, 120, 160, 3)] | 0 |
| conv2d (Conv2D) | (None, 58, 78, 32) | 2432 |
| conv2d_1 (Conv2D) | (None, 27, 37, 32) | 25632 |
| dropout (Dropout) | (None, 27, 37, 32) | 0 |
| conv2d_2 (Conv2D) | (None, 25, 35, 48) | 13872 |
| conv2d_3 (Conv2D) | (None, 23, 33, 48) | 20784 |
| max_pooling2d (MaxPooling2D) | (None, 11, 16, 48) | 0 |
| conv2d_4 (Conv2D) | (None, 9, 14, 64) | 27712 |
| conv2d_5 (Conv2D) | (None, 7, 12, 64) | 36928 |
| max_pooling2d_1 (MaxPooling2D) | (None, 3, 6, 64) | 0 |
| flatten (Flatten) | (None, 1152) | 0 |
| dropout_1 (Dropout) | (None, 1152) | 0 |
| dense (Dense) | (None, 256) | 295168 |
| dense_1 (Dense) | (None, 100) | 25700 |
| batch_normalization (Batch Normalization) | (None, 100) | 400 |
| dense_2 (Dense) | (None, 50) | 5050 |
| dense_3 (Dense) | (None, 10) | 510 |
| dense_4 (Dense) | (None, 2) | 22 |
| Total params: 454,210 | | |
| Trainable params: 454,010 | | |
| Non-trainable params: 200 | | |

The Sweet & Simple Model Predictions on Test Data


$A = 0.92$ $S = 0.99$




$A = 0.45$ $S = 0.99$




$A = 0.61$ $S = 0$



$A = 0.03$ $S = 0.99$



$A = 0.48$ $S = 0$



$A = 0.18$ $S = 0.99$



Thank you for listening

...